

Final Project

December 9, 2019

1 Final Project_STP 598 Topic: Machine Learning / Statistical Learning

Team Members: Chi Duan, Huai Shen and Chengyu hong

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
from itertools import combinations
import matplotlib.pyplot as plt
from nltk.tokenize import RegexpTokenizer
```

```
[2]: from sklearn.ensemble import RandomForestClassifier, BaggingClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.naive_bayes import GaussianNB
from sklearn import tree
from sklearn.svm import SVC
from sklearn.model_selection import cross_validate
from sklearn.metrics import mean_squared_error, accuracy_score, f1_score,
    ↪ confusion_matrix
```

```
[3]: import keras
      from keras.models import Sequential
      from keras.layers import Dense, Dropout, MaxPooling1D, Conv1D
```

Using TensorFlow backend.

```
[4]: # Wine data cleaning
def wine_data_cleaning(wine_data):
    cleaned_data = []
    column_name = wine_data.columns.values[0].split(';')
    for i in range(len(column_name)):
        column_name[i] = column_name[i].strip(' ')
    for i in range(len(wine_data)):
        cleaned_data.append(list(map(float, wine_data.iloc[i].values[0].split(';')
        )))
```

```
cleaned_data = pd.DataFrame(cleaned_data, columns=column_name)
return cleaned_data
```

```
[5]: def merge_quality(y,label_to_merge=[5,6]):
      for i in range(len(y)):
          if y[i] in label_to_merge:
              y[i]=label_to_merge[0]
      return y
```

2 Red Wine Data

```
[6]: # Read the Red wine data
red_wine_raw = pd.read_csv('winequality-red.csv')
red_wine_cleaned = wine_data_cleaning(red_wine_raw)
print('Does each attribute column have null data? \n', red_wine_cleaned.
      ↪isnull().any())
# Show the data statistics
red_wine_cleaned.describe()
```

Does each attribute column have null data?

fixed acidity	False
volatile acidity	False
citric acid	False
residual sugar	False
chlorides	False
free sulfur dioxide	False
total sulfur dioxide	False
density	False
pH	False
sulphates	False
alcohol	False
quality	False

dtype: bool

```
[6]:
```

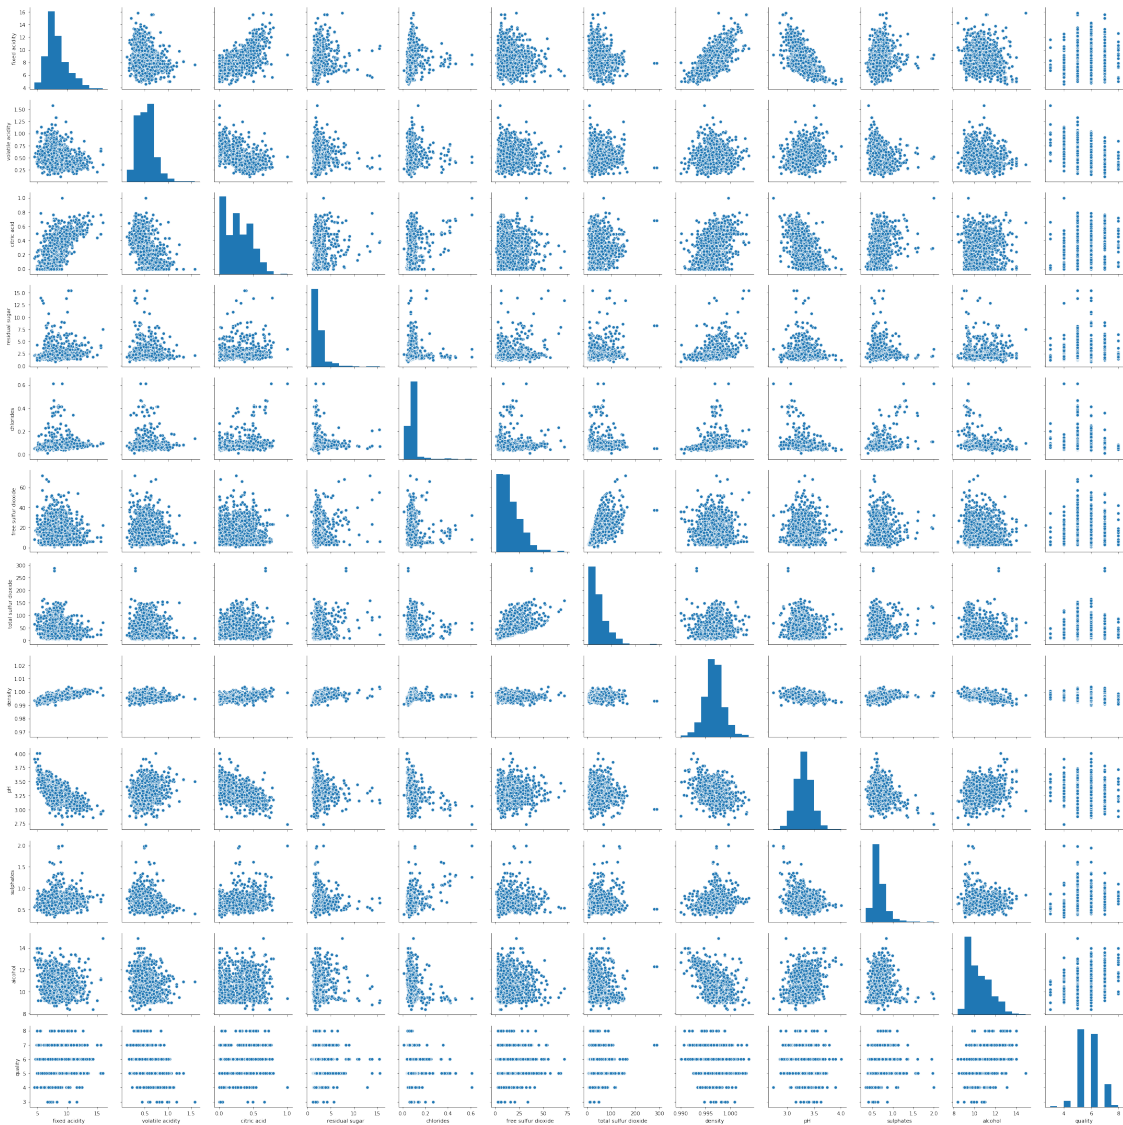
	fixed acidity	volatile acidity	citric acid	residual sugar	\
count	1599.000000	1599.000000	1599.000000	1599.000000	
mean	8.319637	0.527821	0.270976	2.538806	
std	1.741096	0.179060	0.194801	1.409928	
min	4.600000	0.120000	0.000000	0.900000	
25%	7.100000	0.390000	0.090000	1.900000	
50%	7.900000	0.520000	0.260000	2.200000	
75%	9.200000	0.640000	0.420000	2.600000	
max	15.900000	1.580000	1.000000	15.500000	

	chlorides	free sulfur dioxide	total sulfur dioxide	density	\
count	1599.000000	1599.000000	1599.000000	1599.000000	

mean	0.087467	15.874922	46.467792	0.996747
std	0.047065	10.460157	32.895324	0.001887
min	0.012000	1.000000	6.000000	0.990070
25%	0.070000	7.000000	22.000000	0.995600
50%	0.079000	14.000000	38.000000	0.996750
75%	0.090000	21.000000	62.000000	0.997835
max	0.611000	72.000000	289.000000	1.003690


	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	3.311113	0.658149	10.422983	5.636023
std	0.154386	0.169507	1.065668	0.807569
min	2.740000	0.330000	8.400000	3.000000
25%	3.210000	0.550000	9.500000	5.000000
50%	3.310000	0.620000	10.200000	6.000000
75%	3.400000	0.730000	11.100000	6.000000
max	4.010000	2.000000	14.900000	8.000000

```
[7]: # Show the paire-wise attribute correlation
g = sns.pairplot(red_wine_cleaned)
```



```
[8]: # Plot the attribute correlation coefficient as in heat map
plt.figure(figsize=(10, 10))
ax = sns.heatmap(red_wine_cleaned.corr(),annot=True,vmin=-1,cmap='YlGnBu')
ax.set(xlabel='Attribute \n Red Wine', ylabel='Attribute')
```

```
[8]: [Text(69.0, 0.5, 'Attribute'), Text(0.5, 69.0, 'Attribute \n Red Wine')]
```



```
[9]: # Feature selection by using kNN to find the most salient feature related to
      ↪ wine quality
def feature_selection_red(wine_cleaned, n_neighbors=90):
    # Generate all possible combinations of attribute feature as index array
    ↪ for feature selection
    feature_index = []
    for i in range(1, len(wine_cleaned.columns)-1):
        arr = range(len(wine_cleaned.columns)-1)
        combination = list(combinations(arr, i))
        for c in combination:
            feature_index.append(np.array(c))
    score = []
```

```

    for feature in feature_index:
        X = wine_cleaned.iloc[:,feature].values
        y = wine_cleaned.iloc[:, -1].values
        scaler = StandardScaler()
        scaler.fit(X)
        X_scaled = scaler.transform(X)
        X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
→test_size=0.2, random_state=3)
        # By trial, we used neighbors = 90
    #         n_neighbors = 90
        kNN_clf = KNeighborsClassifier(n_neighbors=n_neighbors).fit(X_train,
→y_train)
        y_predicted = kNN_clf.predict(X_test)
        score.append(accuracy_score(y_predicted, y_test))
    return feature_index[np.argmax(score)], max(score)
best_feature_red, feature_score_red = feature_selection_red(red_wine_cleaned,
→n_neighbors=90)
print('The best feature selection score: ', feature_score_red)
print('The best feature combinations: ', best_feature_red)

```

The best feature selection score: 0.646875
The best feature combinations: [1 4 6 9 10]

```

[19]: # Use the best feature combination as training and testing data
X = red_wine_cleaned.iloc[:, best_feature_red]
y = red_wine_cleaned.iloc[:, -1].values

# If one range of labels need to be merged to one label, set the following
# variable merge_label to True. Else set merge_label=False.
merge_label = False
if merge_label:
    label_to_merge = [5, 6]
    y = merge_quality(y, label_to_merge= label_to_merge)

scaler = StandardScaler()
scaler.fit(X)
X_scaled = scaler.transform(X)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
→random_state=3)

```

```

[20]: # Use Gaussian Navie Bayes classifier to model the data
def NB_clf(X_train, X_test, y_train, y_test):
    NB_clf = GaussianNB()
    NB_clf.fit(X_train, y_train)
    y_predicted = NB_clf.predict(X_test)

```

```

    NB_clf_error = [mean_squared_error(y_predicted, y_test),
↳accuracy_score(y_predicted, y_test),f1_score(y_predicted, y_test,
↳average='micro')]
    print('NB_clf')
    print('mean_squared_error: ', mean_squared_error(y_predicted, y_test))
    print('accuracy_score: ', accuracy_score(y_predicted, y_test))
    print('f1_score: ', f1_score(y_predicted, y_test, average='micro'))
    c_matrix = confusion_matrix(y_test, y_predicted)
    print('the prediction confusion matrix is: \n', c_matrix)
    return NB_clf_error
NB_clf_error_red = NB_clf(X_train, X_test, y_train, y_test)

```

```

NB_clf
mean_squared_error:  0.484375
accuracy_score:  0.609375
f1_score:  0.609375
the prediction confusion matrix is:
[[ 0  1  1  0  0  0]
 [ 0  1  7  4  0  0]
 [ 0  7 97 33  1  0]
 [ 0  0 37 81 11  2]
 [ 0  0  1 16 16  3]
 [ 0  0  0  1  0  0]]

```

```

[18]: # Use kNN classifier to model the data and use grid search to find the optimal
↳model parameters
parameter_tree = {'n_neighbors':range(10,100)}
def kNN_clf(X_train, X_test, y_train, y_test, parameter_tree={'n_neighbors':
↳range(10,100)}):
    kNN_clf = KNeighborsClassifier()
    clf = GridSearchCV(kNN_clf, parameter_tree, cv=5, iid=False, n_jobs=8)
    clf.fit(X_train, y_train)
    y_predicted = clf.predict(X_test)
    kNN_clf_error = [mean_squared_error(y_predicted, y_test),
↳accuracy_score(y_predicted, y_test),f1_score(y_predicted, y_test,
↳average='micro')]
    print('kNN_clf')
    print('mean_squared_error: ', mean_squared_error(y_predicted, y_test))
    print('accuracy_score: ', accuracy_score(y_predicted, y_test))
    print('f1_score: ', f1_score(y_predicted, y_test, average='micro'))
    c_matrix = confusion_matrix(y_test, y_predicted)
    print('the prediction confusion matrix is: \n', c_matrix)
    kNN_para = clf.best_params_
    print('The best parameter for kNN_clf: ', kNN_para)
    return kNN_clf_error
kNN_clf_error_red = kNN_clf(X_train, X_test, y_train, y_test,
↳parameter_tree=parameter_tree)

```

```

kNN_clf
mean_squared_error: 0.403125
accuracy_score: 0.653125
f1_score: 0.653125
the prediction confusion matrix is:
[[ 0  0  2  0  0  0]
 [ 0  0  9  3  0  0]
 [ 0  0 104 34  0  0]
 [ 0  0 35 93  3  0]
 [ 0  0  0 24 12  0]
 [ 0  0  0  1  0  0]]
The best parameter for kNN_clf: {'n_neighbors': 88}

```

```

[22]: # Use Decision Tree classifier to model the data and use grid search to find
      ↪ the optimal model parameters
parameter_tree = {'criterion': ['gini', 'entropy'], 'splitter': ['best',
      ↪ 'random'] ,
                  'max_depth': range(1,25), 'min_samples_split': range(2,10),
                  'min_samples_leaf': range(1,10)}
def tree_clf(X_train, X_test, y_train, y_test, parameter_tree={'criterion':
      ↪ ['gini', 'entropy']}):
    tree_clf = tree.DecisionTreeClassifier()
    clf = GridSearchCV(tree_clf, parameter_tree, cv=5, iid=False, n_jobs=8)
    clf.fit(X_train, y_train)
    y_predicted = clf.predict(X_test)
    tree_clf_error_red = [mean_squared_error(y_predicted, y_test),
      ↪ accuracy_score(y_predicted, y_test), f1_score(y_predicted, y_test,
      ↪ average='micro')]
    print('tree_clf')
    print('mean_squared_error: ', mean_squared_error(y_predicted, y_test))
    print('accuracy_score: ', accuracy_score(y_predicted, y_test))
    print('f1_score: ', f1_score(y_predicted, y_test, average='micro'))
    c_matrix = confusion_matrix(y_test, y_predicted)
    print('the prediction confusion matrix is: \n', c_matrix)
    decision_tree_para = clf.best_params_
    print('The best parameter for tree_clf: ', decision_tree_para)
    return tree_clf_error_red
tree_clf_error_red = tree_clf(X_train, X_test, y_train, y_test,
      ↪ parameter_tree=parameter_tree)

```

```

tree_clf
mean_squared_error: 0.6125
accuracy_score: 0.646875
f1_score: 0.646875
the prediction confusion matrix is:
[[ 0  1  1  0  0  0]
 [ 0  2  7  2  1  0]

```



```
[ 1  4 101  27  5  0]
[ 0  3  27  86 12  3]
[ 0  0  9   6 18  3]
[ 0  0  0   1  0  0]]
```

The best parameter for tree_clf: {'criterion': 'entropy', 'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 2, 'splitter': 'best'}

```
[23]: # Use Random Forest classifier to model the data and use grid search to find
      ↳ the optimal model parameters
parameter_tree = {'n_estimators': range(35,95, 2), 'criterion': ['gini',
      ↳ 'entropy'],
                  'max_depth': range(12,30,2), 'min_samples_split': range(2,5,1),
                  'min_samples_leaf': range(1,2,1)}
def random_forest(X_train, X_test, y_train, y_test,
      ↳ parameter_tree={'n_estimators': range(35,85, 15)}):
    random_forest = RandomForestClassifier(verbose=1)
    clf = GridSearchCV(random_forest, parameter_tree, cv=5, iid=False,
      ↳ n_jobs=8, verbose=1)
    clf.fit(X_train, y_train)
    y_predicted = clf.predict(X_test)
    random_forest_error = [mean_squared_error(y_predicted, y_test),
      ↳ accuracy_score(y_predicted, y_test), f1_score(y_predicted, y_test,
      ↳ average='micro')]
    print('random_forest')
    print('mean_squared_error: ', mean_squared_error(y_predicted, y_test))
    print('accuracy_score: ', accuracy_score(y_predicted, y_test))
    print('f1_score: ', f1_score(y_predicted, y_test, average='micro'))
    c_matrix = confusion_matrix(y_test, y_predicted)
    print('the prediction confusion matrix is: \n', c_matrix)
    random_forest_para = clf.best_params_
    print('The best parameter for random_forest: ', random_forest_para)
    return random_forest_error
random_forest_error_red=random_forest(X_train, X_test, y_train, y_test,
      ↳ parameter_tree=parameter_tree)
```

Fitting 5 folds for each of 1620 candidates, totalling 8100 fits

```
[Parallel(n_jobs=8)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done 52 tasks      | elapsed: 0.9s
[Parallel(n_jobs=8)]: Done 352 tasks     | elapsed: 7.4s
[Parallel(n_jobs=8)]: Done 852 tasks     | elapsed: 19.0s
[Parallel(n_jobs=8)]: Done 1552 tasks    | elapsed: 37.6s
[Parallel(n_jobs=8)]: Done 2452 tasks    | elapsed: 1.1min
[Parallel(n_jobs=8)]: Done 3552 tasks    | elapsed: 1.6min
[Parallel(n_jobs=8)]: Done 4528 tasks    | elapsed: 2.3min
[Parallel(n_jobs=8)]: Done 5278 tasks    | elapsed: 3.0min
[Parallel(n_jobs=8)]: Done 6128 tasks    | elapsed: 3.7min
[Parallel(n_jobs=8)]: Done 7078 tasks    | elapsed: 4.5min
```

```
[Parallel(n_jobs=8)]: Done 8100 out of 8100 | elapsed: 5.3min finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 55 out of 55 | elapsed: 0.1s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
```

```
random_forest
mean_squared_error: 0.353125
accuracy_score: 0.721875
f1_score: 0.7218749999999999
the prediction confusion matrix is:
[[ 0  0  2  0  0  0]
 [ 0  0  8  4  0  0]
 [ 0  2 119 17  0  0]
 [ 0  0 27 94 10  0]
 [ 0  0  1 13 18  4]
 [ 0  0  0  1  0  0]]
The best parameter for random_forest: {'criterion': 'entropy', 'max_depth': 18,
'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 55}
[Parallel(n_jobs=1)]: Done 55 out of 55 | elapsed: 0.0s finished
```

```
[28]: #Turn the predicted probability into class labels
def translate_label(y):
    y_translated = []
    for i in y:
        y_translated.append(np.argmax(i))
    return y_translated
```

```
[27]: # Neural Network - Densely Connected
X = red_wine_cleaned.iloc[:, :-1]
y = red_wine_cleaned.iloc[:, -1].values

scaler = StandardScaler()
scaler.fit(X)
X_scaled = scaler.transform(X)
# print(X_scaled)
num_classes = 10
y = keras.utils.to_categorical(y, num_classes)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
    ↪random_state=7)
d = 0.15
model = Sequential()
model.add(Dense(30, input_dim=len(X_train[0]), activation='linear'))
model.add(Dropout(d))
model.add(Dense(50, activation='tanh'))
model.add(Dropout(d))
model.add(Dense(50, activation='tanh'))
model.add(Dropout(d))
```

```

model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
optimizer=keras.optimizers.Adadelta(),
metrics=['accuracy'])

batch_size = 50
model.fit(X_train, y_train, batch_size=batch_size, epochs=300, verbose=0,
→validation_data=(X_test, y_test))
y_predicted = model.predict(X_test)

```

```

[26]: y_predicted_scalar = translate_label(y_predicted)
y_test_scalar = translate_label(y_test)
neural_network_red = [mean_squared_error(y_predicted_scalar, y_test_scalar),
accuracy_score(y_predicted_scalar, y_test_scalar),
f1_score(y_predicted_scalar, y_test_scalar,
→average='micro')]
print('densely connected neuralwork')
print('mean_squared_error: ', mean_squared_error(y_predicted_scalar,
→y_test_scalar))
print('accuracy_score: ', accuracy_score(y_predicted_scalar, y_test_scalar))
print('f1_score: ', f1_score(y_predicted_scalar, y_test_scalar,
→average='micro'))
c_matrix = confusion_matrix(y_predicted_scalar, y_test_scalar)
print('the prediction confusion matrix is: \n', c_matrix)

```

```

densely connected neuralwork
mean_squared_error: 0.40625
accuracy_score: 0.675
f1_score: 0.675
the prediction confusion matrix is:
[[ 1  1  0  0  0]
 [ 7 95 31  0  0]
 [ 1 25 101 16  2]
 [ 1  2 13 19  2]
 [ 0  0  1  2  0]]

```

3 White Wine Data

```

[29]: # Read the Red wine data
white_wine_raw = pd.read_csv('winequality-white.csv')
white_wine_cleaned = wine_data_cleaning(white_wine_raw)
# We fitered out wine with quality=9 since there are only 5 instance for this
→category, which is
# too less for a classifier
white_wine_cleaned = white_wine_cleaned[white_wine_cleaned['quality']!=9]

```

```
print('Does each attribute column have null data? \n', red_wine_cleaned.
      ↪isnull().any())
# Show the data statistics
white_wine_cleaned.describe()
```

Does each attribute column have null data?

```
fixed acidity      False
volatile acidity   False
citric acid        False
residual sugar     False
chlorides          False
free sulfur dioxide False
total sulfur dioxide False
density            False
pH                 False
sulphates          False
alcohol            False
quality            False
dtype: bool
```

```
[29]:
```

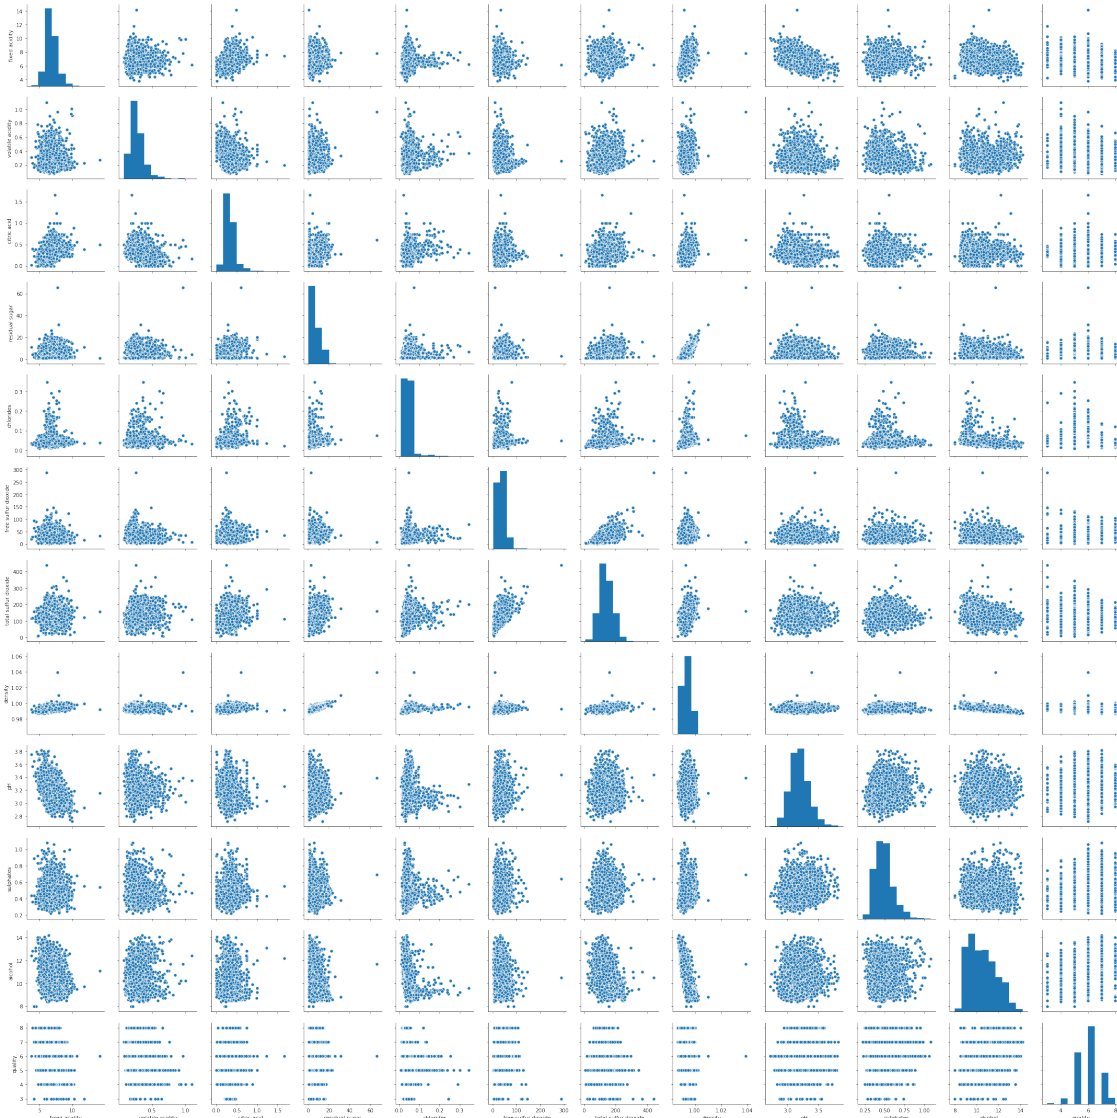
	fixed acidity	volatile acidity	citric acid	residual sugar	\
count	4893.000000	4893.000000	4893.000000	4893.000000	
mean	6.854210	0.278221	0.334139	6.393736	
std	0.843637	0.100831	0.121048	5.072990	
min	3.800000	0.080000	0.000000	0.600000	
25%	6.300000	0.210000	0.270000	1.700000	
50%	6.800000	0.260000	0.320000	5.200000	
75%	7.300000	0.320000	0.390000	9.900000	
max	14.200000	1.100000	1.660000	65.800000	

	chlorides	free sulfur dioxide	total sulfur dioxide	density	\
count	4893.000000	4893.000000	4893.000000	4893.000000	
mean	0.045791	35.310035	138.383507	0.99403	
std	0.021850	17.011384	42.509982	0.00299	
min	0.009000	2.000000	9.000000	0.98711	
25%	0.036000	23.000000	108.000000	0.99173	
50%	0.043000	34.000000	134.000000	0.99375	
75%	0.050000	46.000000	167.000000	0.99610	
max	0.346000	289.000000	440.000000	1.03898	

	pH	sulphates	alcohol	quality
count	4893.000000	4893.000000	4893.000000	4893.000000
mean	3.188144	0.489871	10.512565	5.874719
std	0.151011	0.114151	1.229755	0.880446
min	2.720000	0.220000	8.000000	3.000000
25%	3.090000	0.410000	9.500000	5.000000
50%	3.180000	0.470000	10.400000	6.000000


75%	3.280000	0.550000	11.400000	6.000000
max	3.820000	1.080000	14.200000	8.000000

```
[30]: # Show the paire-wise attribute correlation
g = sns.pairplot(white_wine_cleaned)
```



```
[31]: # Plot the attribute correlation coefficient as in heat map
plt.figure(figsize=(10, 10))
ax = sns.heatmap(white_wine_cleaned.corr(),annot=True,vmin=-1,cmap='YlGnBu')
ax.set(xlabel='Attribute \n White Wine', ylabel='Attribute')
```

```
[31]: [Text(69.0, 0.5, 'Attribute'), Text(0.5, 69.0, 'Attribute \n White Wine')]
```



```
[32]: # Feature selection by using kNN to find the most salient feature related to
      ↪ wine quality
def feature_selection_white(wine_cleaned):
    # Generate all possible combinations of attribute feature as index array
    ↪ for feature selection
    feature_index = []
    for i in range(1, len(wine_cleaned.columns)-1):
        arr = range(len(wine_cleaned.columns)-1)
        combination = list(combinations(arr, i))
        for c in combination:
            feature_index.append(np.array(c))
    score = []
```

```

for feature in feature_index:
    X = wine_cleaned.iloc[:,feature].values
    y = wine_cleaned.iloc[:, -1].values
    scaler = StandardScaler()
    scaler.fit(X)
    X_scaled = scaler.transform(X)
    X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
→test_size=0.2, random_state=3)
    # By trial, we used neighbors = 90
    n_neighbors=90
    kNN_clf = KNeighborsClassifier(n_neighbors=n_neighbors).fit(X_train,
→y_train)
#     NB_clf = GaussianNB().fit(X_train, y_train)
    y_predicted = kNN_clf.predict(X_test)
    score.append(accuracy_score(y_predicted, y_test))
    print('Total number of feature combinations: ', len(score))
    return feature_index[np.argmax(score)], max(score)
best_feature_white, feature_score_white =
→feature_selection_white(white_wine_cleaned)
print('The best feature selection score: ', feature_score_white)
print('The best feature combinations: ', best_feature_white)

```

Total number of feature combinations: 2046

The best feature selection score: 0.5556690500510725

The best feature combinations: [0 1 2 4 5 8 10]

```

[33]: # Use the best feature combination as training and testing data
# print(best_feature_white)
X = white_wine_cleaned.iloc[:,best_feature_white]
y = white_wine_cleaned.iloc[:, -1].values

# If one range of labels need to be merged to one label, set the following
# variable merge_label to True. Else set merge_label=False.
merge_label = False
if merge_label:
    label_to_merge = [5, 6]
    y = merge_quality(y, label_to_merge= label_to_merge)

scaler = StandardScaler()
scaler.fit(X)
X_scaled = scaler.transform(X)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
→random_state=3)
NB_clf = GaussianNB()
NB_clf.fit(X_train, y_train)
y_predicted = NB_clf.predict(X_test)
print('mean_squared_error: ', mean_squared_error(y_predicted, y_test))

```

```
print('accuracy_score: ', accuracy_score(y_predicted, y_test))
print('f1_score: ', f1_score(y_predicted, y_test, average='micro'))
```

```
mean_squared_error: 0.8784473953013279
accuracy_score: 0.48314606741573035
f1_score: 0.48314606741573035
```

```
[52]: parameter_tree = {'n_neighbors':range(10,300,5)}
      kNN_clf_error_red = kNN_clf(X_train, X_test, y_train, y_test,
      ↪parameter_tree=parameter_tree)
```

```
kNN_clf
mean_squared_error: 0.634320735444331
accuracy_score: 0.550561797752809
f1_score: 0.550561797752809
the prediction confusion matrix is:
[[ 0  0  2  0  0  0]
 [ 0  2 20 13  0  0]
 [ 0  1 85 108  7  0]
 [ 0  0 93 302 37  0]
 [ 0  0  8 112 50  0]
 [ 0  0  2 25 12  0]]
The best parameter for kNN_clf: {'n_neighbors': 25}
```

```
[53]: # Use Decision Tree classifier to model the data and use search to find the
      ↪optimal model parameters
      parameter_tree ={'criterion': ['gini', 'entropy'], 'splitter': ['best',
      ↪'random'] ,
                        'max_depth': range(1,35), 'min_samples_split': range(2,10),
                        'min_samples_leaf': range(1,10)}
      tree_clf_error_white = tree_clf(X_train, X_test, y_train, y_test,
      ↪parameter_tree=parameter_tree)
```

```
tree_clf
mean_squared_error: 0.8580183861082737
accuracy_score: 0.5985699693564862
f1_score: 0.5985699693564862
the prediction confusion matrix is:
[[ 0  0  1  1  0  0]
 [ 1  5 10 17  2  0]
 [ 2  9 85 70 17  7]
 [ 4  9 81 278 51  9]
 [ 2  2 14 51 94  7]
 [ 0  2  2 12 10 13]]
The best parameter for tree_clf: {'criterion': 'gini', 'max_depth': 23,
'min_samples_leaf': 1, 'min_samples_split': 2, 'splitter': 'random'}
```



```
[54]: # Use Random Forest classifier to model the data and use grid search to find
      ↪ the optimal model parameters
parameter_tree = {'n_estimators': range(330,380, 5), 'criterion': ['gini',
      ↪ 'entropy'],
                  'max_depth': range(12,20,2), 'min_samples_split': range(2,6,1),
                  'min_samples_leaf': range(1,2,1)}

random_forest = RandomForestClassifier(verbose=1)
clf = GridSearchCV(random_forest, parameter_tree, cv=5, iid=False, n_jobs=8,
      ↪ verbose=1)
clf.fit(X_train, y_train)
y_predicted = clf.predict(X_test)
random_forest_red = [mean_squared_error(y_predicted, y_test),
      ↪ accuracy_score(y_predicted, y_test), f1_score(y_predicted, y_test,
      ↪ average='micro')]
print('random_forest')
print('mean_squared_error: ', mean_squared_error(y_predicted, y_test))
print('accuracy_score: ', accuracy_score(y_predicted, y_test))
print('f1_score: ', f1_score(y_predicted, y_test, average='micro'))
c_matrix = confusion_matrix(y_test, y_predicted)
print('the prediction confusion matrix is: \n', c_matrix)
random_forest_para = clf.best_params_
print('The best parameter for random_forest: ', random_forest_para)
```

Fitting 5 folds for each of 320 candidates, totalling 1600 fits

```
[Parallel(n_jobs=8)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed: 12.4s
[Parallel(n_jobs=8)]: Done 184 tasks     | elapsed: 1.0min
[Parallel(n_jobs=8)]: Done 434 tasks     | elapsed: 2.6min
[Parallel(n_jobs=8)]: Done 784 tasks     | elapsed: 4.8min
[Parallel(n_jobs=8)]: Done 1234 tasks    | elapsed: 8.6min
[Parallel(n_jobs=8)]: Done 1600 out of 1600 | elapsed: 12.1min finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
```

```
random_forest
mean_squared_error: 0.45658835546475995
accuracy_score: 0.6710929519918284
f1_score: 0.6710929519918284
the prediction confusion matrix is:
```

```
[[ 0  0  2  0  0  0]
 [ 0  8 16 11  0  0]
 [ 0  2 194 101  3  1]
 [ 0  0 64 346 22  0]
 [ 0  0  3 69 97  1]
 [ 0  0  0 20  7 12]]
```

```
The best parameter for random_forest: {'criterion': 'entropy', 'max_depth': 18,
'min_samples_leaf': 1, 'min_samples_split': 4, 'n_estimators': 345}
```

```
[Parallel(n_jobs=1)]: Done 345 out of 345 | elapsed: 2.0s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 345 out of 345 | elapsed: 0.0s finished
```

```
[36]: # Neural Network - Densely Connected
X = white_wine_cleaned.iloc[:,best_feature_white]
y = white_wine_cleaned.iloc[:, -1].values
print(best_feature_white)
scaler = StandardScaler()
scaler.fit(X)
X_scaled = scaler.transform(X)
num_classes = 10
y = keras.utils.to_categorical(y, num_classes)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
    ↪random_state=7)
d = 0.05
model = Sequential()
model.add(Dense(50, input_dim=len(X_train[0]), activation='linear'))
model.add(Dropout(d))
model.add(Dense(80, activation='tanh'))
model.add(Dropout(d))
model.add(Dense(50, activation='tanh'))
model.add(Dropout(d))
model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
    optimizer=keras.optimizers.Adadelta(),
    metrics=['accuracy'])

batch_size = 50
model.fit(X_train, y_train, batch_size=batch_size, epochs=300, verbose=0,
    ↪validation_data=(X_test, y_test))
y_predicted = model.predict(X_test)
```

```
[ 0  1  2  4  5  8 10]
```

```
[37]: y_predicted = translate_label(y_predicted)
y_test = translate_label(y_test)
neural_network_red = [mean_squared_error(y_predicted, y_test),
    ↪accuracy_score(y_predicted, y_test),f1_score(y_predicted, y_test,
    ↪average='micro')]
print('densely connected neuralwork')
print('mean_squared_error: ', mean_squared_error(y_predicted, y_test))
print('accuracy_score: ', accuracy_score(y_predicted, y_test))
print('f1_score: ', f1_score(y_predicted, y_test, average='micro'))
c_matrix = confusion_matrix(y_test, y_predicted)
print('the prediction confusion matrix is: \n', c_matrix)
```

```
densely connected neuralwork
```

mean_squared_error: 0.5413687436159347

accuracy_score: 0.6016343207354443

f1_score: 0.6016343207354443

the prediction confusion matrix is:

```
[[ 1  0  3  0  0  0]
```

```
[ 0 12  9  6  0  0]
```

```
[ 3 11 187 87  7  0]
```

```
[ 0  2 85 291 58  2]
```

```
[ 0  0  9 72 85  8]
```

```
[ 0  0  1 12 15 13]]
```

[]:

[]: